

• • •

Operations Research & Optimization

2026 7 31

- 1 (Linear Programming)
- 2 (Nonlinear Programming)
- 3 (Heuristics & Metaheuristics)
- 4 (Convex Optimization)
- 5 ML
- 6 GPU / CUDA ML

(LP)


- 
- /
- LP

(Standard Form)

$$\begin{array}{ll}\min_{\mathbf{x}} & \mathbf{c}^\top \mathbf{x} \\ \text{s.t.} & \mathbf{Ax} = \mathbf{b} \\ & \mathbf{x} \geq \mathbf{0}\end{array}$$

$$\mathbf{A} \in \mathbb{R}^{m \times n} \quad \mathbf{b} \in \mathbb{R}^m \quad \mathbf{c} \in \mathbb{R}^n$$

- **(Polyhedron)**  $\{x \mid Ax = b, x \geq 0\}$
- (Convex Polytope)
- **(extreme point / basic feasible solution)**
- (Simplex Method)



LP

- $\max \mathbf{c}^\top \mathbf{x} \iff \min -\mathbf{c}^\top \mathbf{x}$
- $\mathbf{Ax} \leq \mathbf{b} \iff \mathbf{Ax} + \mathbf{s} = \mathbf{b}, \mathbf{s} \geq \mathbf{0}$
- $\mathbf{Ax} \geq \mathbf{b} \iff \mathbf{Ax} - \mathbf{s} = \mathbf{b}, \mathbf{s} \geq \mathbf{0}$
- $x_i \quad \iff \quad x_i = x_i^+ - x_i^-, x_i^+, x_i^- \geq 0$

LP

/

Simplex

(Basic Feasible Solution, BFS)  $A \quad m \quad B \quad 0 \quad Bx_B = b$

- 1 BFS
- 2 (reduced cost)  $\bar{c}_N^\top = c_N^\top - c_B^\top B^{-1}N$
- 3  $\bar{c}_N \geq 0$
- 4 Dantzig (ratio test)
- 5

# (Simplex Tableau)

LP

|                | $\mathbf{x}_B$ | $\mathbf{x}_N$              | RHS                                            |
|----------------|----------------|-----------------------------|------------------------------------------------|
| $\mathbf{x}_B$ | $\mathbf{I}$   | $\mathbf{B}^{-1}\mathbf{N}$ | $\mathbf{B}^{-1}\mathbf{b}$                    |
| $z$            | $\mathbf{0}$   | $\bar{\mathbf{c}}_N^\top$   | $\mathbf{c}_B^\top \mathbf{B}^{-1} \mathbf{b}$ |

- $\bar{\mathbf{c}}_N \geq \mathbf{0}$
- $\bar{c}_j < 0 \leq \mathbf{0}$
- Min ratio test

# (Two-Phase Method)

BFS

■ **Phase I**       $\mathbf{a}$

$$\min \sum_i a_i \quad \text{s.t.} \quad \mathbf{Ax} + \mathbf{a} = \mathbf{b}, \quad \mathbf{x}, \mathbf{a} \geq \mathbf{0}$$

$> 0$

■ **Phase II**    Phase I    BFS

■  $M$        $M \sum a_i$



(Primal)

$$\min_{\mathbf{x}} \mathbf{c}^\top \mathbf{x} \quad \text{s.t.} \quad \mathbf{Ax} = \mathbf{b}, \mathbf{x} \geq \mathbf{0}$$

(Dual)

$$\max_{\mathbf{y}} \mathbf{b}^\top \mathbf{y} \quad \text{s.t.} \quad \mathbf{A}^\top \mathbf{y} \leq \mathbf{c}$$

| Primal ( )                              | Dual ( )                                |
|-----------------------------------------|-----------------------------------------|
| $x_i \geq 0$                            | $\mathbf{a}_i^\top \mathbf{y} \leq c_i$ |
| $x_i$                                   | $\mathbf{a}_i^\top \mathbf{y} = c_i$    |
| $\mathbf{a}_i^\top \mathbf{x} \geq b_i$ | $y_i \geq 0$                            |
| $\mathbf{a}_i^\top \mathbf{x} = b_i$    | $y_i$                                   |

### (Weak Duality)

$\mathbf{x}$  Primal    $\mathbf{y}$  Dual

$$\mathbf{c}^\top \mathbf{x} \geq \mathbf{b}^\top \mathbf{y}$$

### (Strong Duality)

Primal          Dual

$$\min \mathbf{c}^\top \mathbf{x} = \max \mathbf{b}^\top \mathbf{y}$$

### (Complementary Slackness)

$$x_i^* \cdot (\mathbf{a}_i^\top \mathbf{y}^* - c_i) = 0 \quad i$$

$$y_i^* \cdot (\mathbf{a}_i^\top \mathbf{x}^* - b_i) = 0 \quad i$$

- $y$  (Shadow Price)

- $y_i = \frac{\partial \text{Optimal Value}}{\partial b_i}$

- $y_i \geq 0$

- $y_i > 0$

(Sensitivity Analysis)  $b_i, c_j$

- $B^{-1}b \geq 0$

- $\bar{c}_N \geq 0$

LP      (Interior Point Method)

-      KKT

$$\begin{aligned} \mathbf{Ax} &= \mathbf{b}, \quad \mathbf{x} > \mathbf{0} \\ \mathbf{A}^\top \mathbf{y} + \mathbf{s} &= \mathbf{c}, \quad \mathbf{s} > \mathbf{0} \\ \mathbf{XSe} &= \mu \mathbf{e} \end{aligned}$$

$$\mathbf{X} = \text{diag}(\mathbf{x}), \quad \mathbf{S} = \text{diag}(\mathbf{s}) \quad \mu \rightarrow 0$$

Newton       $\mu$       LP

- 
- $m \quad n$
- Hungarian  $n \quad n$
- $/ \rightarrow$  LP
- 
- **(DEA)**

LP    Gurobi, CPLEX, HiGHS, GLPK

## NLP

$$\begin{array}{ll}\min_{\mathbf{x}} & f(\mathbf{x}) \\ \text{s.t.} & g_i(\mathbf{x}) \leq 0, \quad i = 1, \dots, m \\ & h_j(\mathbf{x}) = 0, \quad j = 1, \dots, p\end{array}$$

$f, g_i, h_j$

- **NLP**  $\min f(\mathbf{x})$
- **NLP**
- $\rightarrow$

$$\min f(\mathbf{x}), \mathbf{x} \in \mathbb{R}^n$$

$$\mathbf{x}^* \quad f \quad \nabla f(\mathbf{x}^*) = \mathbf{0}$$

$$\mathbf{x}^* \quad f \quad \nabla f(\mathbf{x}^*) = \mathbf{0} \quad \nabla^2 f(\mathbf{x}^*) \succeq \mathbf{0}$$

$$\nabla f(\mathbf{x}^*) = \mathbf{0} \quad \nabla^2 f(\mathbf{x}^*) \succ \mathbf{0} \quad \mathbf{x}^*$$

## (Gradient Descent)

$$\mathbf{x}_{k+1} = \mathbf{x}_k - \alpha_k \nabla f(\mathbf{x}_k)$$

- $\alpha_k$  (step size / learning rate)
- $\mathcal{O}(1/k)$
- **Wolfe**
- **Backtracking line search**

ill-conditioned



# (Newton's Method)

Taylor

$$f(\mathbf{x}_k + \mathbf{p}) \approx f(\mathbf{x}_k) + \nabla f(\mathbf{x}_k)^\top \mathbf{p} + \frac{1}{2} \mathbf{p}^\top \nabla^2 f(\mathbf{x}_k) \mathbf{p}$$

Newton  $\mathbf{p}_k = -[\nabla^2 f(\mathbf{x}_k)]^{-1} \nabla f(\mathbf{x}_k)$

$$\mathbf{x}_{k+1} = \mathbf{x}_k + \alpha_k \mathbf{p}_k$$



Hessian  $\mathcal{O}(n^3)$  Hessian

# (Quasi-Newton Methods)

Hessian      Hessian

## BFGS (Broyden–Fletcher–Goldfarb–Shanno)

$$s_k = \mathbf{x}_{k+1} - \mathbf{x}_k$$

$$y_k = \nabla f(\mathbf{x}_{k+1}) - \nabla f(\mathbf{x}_k)$$

$$\mathbf{H}_{k+1} = \left( \mathbf{I} - \rho_k s_k y_k^\top \right) \mathbf{H}_k \left( \mathbf{I} - \rho_k y_k s_k^\top \right) + \rho_k s_k s_k^\top$$

$$\rho_k = 1 / (y_k^\top s_k)$$



■ **L-BFGS**       $m$      $(s_k, y_k)$      $\mathcal{O}(mn)$

# Karush–Kuhn–Tucker (KKT)

NLP — KKT

$$\nabla f(\mathbf{x}^*) + \sum_{i=1}^m \mu_i \nabla g_i(\mathbf{x}^*) + \sum_{j=1}^p \lambda_j \nabla h_j(\mathbf{x}^*) = \mathbf{0}$$

$$g_i(\mathbf{x}^*) \leq 0, \quad i = 1, \dots, m$$

$$h_j(\mathbf{x}^*) = 0, \quad j = 1, \dots, p$$

$$\mu_i \geq 0, \quad i = 1, \dots, m$$

$$\mu_i \cdot g_i(\mathbf{x}^*) = 0, \quad i = 1, \dots, m \quad ( \quad )$$

- $\mu_i$  Lagrange
- $\lambda_j$  Lagrange
- KKT

- **(Penalty Method)**

$$\min_{\mathbf{x}} f(\mathbf{x}) + \rho \left[ \sum_i \max(0, g_i(\mathbf{x}))^2 + \sum_j h_j(\mathbf{x})^2 \right]$$

$\rho \rightarrow \infty$       ill-conditioning

- **(Barrier Method)**

$$\min_{\mathbf{x}} f(\mathbf{x}) - \mu \sum_i \ln(-g_i(\mathbf{x}))$$

$\mu \rightarrow 0$

- **SQP (Sequential Quadratic Programming)**      (QP)

- **Lagrange (Augmented Lagrangian)**

## (Trust Region Methods)

$\Delta_k$

$$\min_{\|\mathbf{p}\| \leq \Delta_k} m_k(\mathbf{p}) = f(\mathbf{x}_k) + \nabla f(\mathbf{x}_k)^\top \mathbf{p} + \frac{1}{2} \mathbf{p}^\top \mathbf{B}_k \mathbf{p}$$

$\rho_k$

$\Delta_k$

$$\rho_k = \frac{f(\mathbf{x}_k) - f(\mathbf{x}_k + \mathbf{p}_k)}{m_k(\mathbf{0}) - m_k(\mathbf{p}_k)}$$



ill-conditioned



Newton

NLP      NP-hard

(Heuristic)

(Metaheuristic)

exploitation ( )    exploration ( )

- |         |         |
|---------|---------|
| ■ /     | ■ (GA)  |
| ■ (SA)  | ■ (PSO) |
| ■ (TS)  | ■ (ACO) |
| ■ (VNS) | ■ (DE)  |

(Greedy)

(Local Search) / (Hill Climbing)

- $N(x)$
- $f(x') < f(x) \quad x'$
- —

“ ”

“ ”

→ → →

## Metropolis

- $\Delta E = f(\mathbf{x}_{\text{new}}) - f(\mathbf{x}) < 0$
  - $\Delta E \geq 0 \quad p = \exp(-\Delta E/T)$
- $T$  cooling schedule  $T$

- $T_0$
- $\alpha \quad T_{k+1} = \alpha \cdot T_k, \quad \alpha \in (0.9, 0.99)$
- $T < T_{\min}$



# (GA)

- 1 (population)
  - 2  $f(\mathbf{x})$
  - 3 **(Selection)** / —
  - 4 **(Crossover)**
  - 5 **(Mutation)**
  - 6 2–5
- TSP
- $p_c \approx 0.7\text{--}0.9$      $p_m \approx 0.001\text{--}0.1$
  - (Elitism)

# (PSO)

$i$

- $\mathbf{x}_i$
- $\mathbf{v}_i$
- $\mathbf{p}_{\text{best},i}$
- $\mathbf{g}_{\text{best}}$

$$\mathbf{v}_i \leftarrow \omega \mathbf{v}_i + c_1 r_1 (\mathbf{p}_{\text{best},i} - \mathbf{x}_i) + c_2 r_2 (\mathbf{g}_{\text{best}} - \mathbf{x}_i)$$

$$\mathbf{x}_i \leftarrow \mathbf{x}_i + \mathbf{v}_i$$

- $\omega$       VS

- $c_1, c_2$     /       $\approx 2$

-

# (Tabu Search)

## (Tabu List)

- “ ”
- (Tabu Tenure)  $\ell$
- (Aspiration Criterion)
- (Long-term Memory)

TSP

# (Ant Colony Optimization, ACO)

(pheromone)

## TSP    ACO

- $\tau_{ij}$      $(i, j)$
- $k$      $i$      $j$

$$p_{ij}^k = \frac{\tau_{ij}^\alpha \cdot \eta_{ij}^\beta}{\sum_{\ell \in N_i^k} \tau_{i\ell}^\alpha \cdot \eta_{i\ell}^\beta}$$

$$\eta_{ij} = 1/d_{ij} \quad /$$

- (Evaporation)  $\tau_{ij} \leftarrow (1 - \rho)\tau_{ij}$

-

| SA  | + | $p_{\text{best}}, g_{\text{best}}$ | / |
|-----|---|------------------------------------|---|
| GA  | + |                                    | / |
| PSO | + |                                    |   |
| TS  | + |                                    |   |
| ACO | + |                                    | / |

**No Free Lunch Theorem**

# (Convex Sets)

$$C \subseteq \mathbb{R}^n \quad \mathbf{x}, \mathbf{y} \in C \quad \theta \in [0, 1]$$

$$\theta \mathbf{x} + (1 - \theta) \mathbf{y} \in C$$

- $\{\mathbf{x} \mid \mathbf{a}^\top \mathbf{x} = b\}$

- $\{\mathbf{x} \mid \mathbf{a}^\top \mathbf{x} \leq b\}$

- $\{\mathbf{x} \mid \mathbf{A}\mathbf{x} \preceq \mathbf{b}\}$

- 

- $\{(\mathbf{x}, t) \mid \|\mathbf{x}\| \leq t\}$

- $\bigcap_{\alpha} C_{\alpha}$

- $f(C)$

- $/$

# (Convex Functions)

$$f: \mathbb{R}^n \rightarrow \mathbb{R}$$

$$\mathbf{x}, \mathbf{y} \quad \theta \in [0, 1]$$

$$f(\theta \mathbf{x} + (1 - \theta) \mathbf{y}) \leq \theta f(\mathbf{x}) + (1 - \theta) f(\mathbf{y})$$

- $f(\mathbf{y}) \geq f(\mathbf{x}) + \nabla f(\mathbf{x})^\top (\mathbf{y} - \mathbf{x})$
  - $\nabla^2 f(\mathbf{x}) \succeq \mathbf{0}$  Hessian
  - (Epigraph)  $\text{epi}(f) = \{(\mathbf{x}, t) \mid f(\mathbf{x}) \leq t\} \iff f$
- $$f(\mathbf{y}) \geq f(\mathbf{x}) + \nabla f(\mathbf{x})^\top (\mathbf{y} - \mathbf{x}) + \frac{\mu}{2} \|\mathbf{y} - \mathbf{x}\|^2 \quad \mu > 0$$

$$\begin{aligned}
 \min_{\mathbf{x}} \quad & f_0(\mathbf{x}) \\
 \text{s.t.} \quad & f_i(\mathbf{x}) \leq 0, \quad i = 1, \dots, m \\
 & \mathbf{Ax} = \mathbf{b}
 \end{aligned}$$

$f_0, f_1, \dots, f_m$

= KKT

- LP, QP ( ), QCQP ( )
- SOCP ( ), SDP ( )
- (Geometric Programming)



Lagrange

$$g(\boldsymbol{\lambda}, \boldsymbol{\nu}) = \inf_{\mathbf{x}} \mathcal{L}(\mathbf{x}, \boldsymbol{\lambda}, \boldsymbol{\nu})$$

$$\mathcal{L}(\mathbf{x}, \boldsymbol{\lambda}, \boldsymbol{\nu}) = f_0(\mathbf{x}) + \sum_i \lambda_i f_i(\mathbf{x}) + \sum_j \nu_j (\mathbf{a}_j^\top \mathbf{x} - b_j)$$

Slater (Slater's Condition)

Primal = Dual

## (Subgradient Method)

$f$

$$\mathbf{x}_{k+1} = \mathbf{x}_k - \alpha_k \mathbf{g}_k, \quad \mathbf{g}_k \in \partial f(\mathbf{x}_k)$$

$\partial f(\mathbf{x}) =$

- $\alpha_k \rightarrow 0, \sum \alpha_k = \infty$
- $\mathcal{O}(1/\sqrt{k})$
- $\ell_1$  SVM

# (Proximal Gradient Method)

$$\min_{\mathbf{x}} f(\mathbf{x}) + g(\mathbf{x}) \quad f \text{ diff} \quad g \text{ prox}$$

**(Proximal Operator)**

$$\text{prox}_{\eta g}(\mathbf{v}) = \arg \min_{\mathbf{x}} \left\{ g(\mathbf{x}) + \frac{1}{2\eta} \|\mathbf{x} - \mathbf{v}\|^2 \right\}$$

**ISTA / FISTA** /

$$\mathbf{z}_k = \mathbf{x}_k - \eta \nabla f(\mathbf{x}_k)$$

$$\mathbf{x}_{k+1} = \text{prox}_{\eta g}(\mathbf{z}_k)$$

- $\ell_1$  (LASSO)
- FISTA  $\mathcal{O}(1/k^2)$

# Nesterov (Accelerated Gradient)

$\mathcal{O}(1/k)$  Nesterov  $\mathcal{O}(1/k^2)$

**Nesterov (NAG)**

$$\mathbf{y}_k = \mathbf{x}_k + \beta_k(\mathbf{x}_k - \mathbf{x}_{k-1})$$

$$\mathbf{x}_{k+1} = \mathbf{y}_k - \alpha_k \nabla f(\mathbf{y}_k)$$

$$\beta_k = \frac{k-1}{k+2}$$

- "Look ahead"

- $\mathcal{O}((1 - \sqrt{\mu/L})^k)$

- Nesterov SGD with Momentum Adam

# ADMM (Alternating Direction Method of Multipliers)

$$\begin{aligned} \min_{\mathbf{x}, \mathbf{z}} \quad & f(\mathbf{x}) + g(\mathbf{z}) \\ \text{s.t.} \quad & \mathbf{Ax} + \mathbf{Bz} = \mathbf{c} \end{aligned}$$

$$\text{Lagrange } \mathcal{L}_\rho(\mathbf{x}, \mathbf{z}, \mathbf{y}) = f(\mathbf{x}) + g(\mathbf{z}) + \mathbf{y}^\top (\mathbf{Ax} + \mathbf{Bz} - \mathbf{c}) + \frac{\rho}{2} \|\mathbf{Ax} + \mathbf{Bz} - \mathbf{c}\|^2$$

ADMM

$$\mathbf{x}_{k+1} = \arg \min_{\mathbf{x}} \mathcal{L}_\rho(\mathbf{x}, \mathbf{z}_k, \mathbf{y}_k)$$

$$\mathbf{z}_{k+1} = \arg \min_{\mathbf{z}} \mathcal{L}_\rho(\mathbf{x}_{k+1}, \mathbf{z}, \mathbf{y}_k)$$

$$\mathbf{y}_{k+1} = \mathbf{y}_k + \rho(\mathbf{Ax}_{k+1} + \mathbf{Bz}_{k+1} - \mathbf{c})$$

# (SDP)

SDP   LP

$$\begin{aligned} \min_{\mathbf{X}} \quad & \langle \mathbf{C}, \mathbf{X} \rangle = \text{tr}(\mathbf{C}^\top \mathbf{X}) \\ \text{s.t.} \quad & \langle \mathbf{A}_i, \mathbf{X} \rangle = b_i, \quad i = 1, \dots, m \\ & \mathbf{X} \succeq \mathbf{0} \quad ( \quad ) \end{aligned}$$

- (Primal-Dual IPM)
- PCA   SDP
- Max-Cut   Goemans–Williamson 0.878   SDP

LP  $\subset$  SOCP  $\subset$  SDP

- CVXPY (Python)
- CVX (MATLAB)
- Convex.jl (Julia)

- LP  $\rightarrow$  LP
- QP  $\rightarrow$  QP
- SOCP/SDP  $\rightarrow$
- NLP (IPOPT, SNOPT )

- ECOS, SCS —SOCP, SDP
- OSQP —QP
- MOSEK — LP/QP/SOCP/SDP
- Gurobi / CPLEX —LP, MILP, QP

$$\min_{\theta} \frac{1}{N} \sum_{i=1}^N \ell(\mathbf{x}_i, y_i; \theta) + \lambda R(\theta)$$

- $\ell(\cdot)$       MSE Hinge
- $R(\cdot)$      $\ell_2$ ,  $\ell_1$ , dropout
- $N$         ML

## ML

- $N \rightarrow$
- $n \rightarrow$
- 
- $\gg$



# (SGD)

(mini-batch)  $\mathcal{B}$

$$\boldsymbol{\theta}_{t+1} = \boldsymbol{\theta}_t - \eta_t \cdot \frac{1}{|\mathcal{B}|} \sum_{i \in \mathcal{B}} \nabla_{\boldsymbol{\theta}} \ell(\mathbf{x}_i, y_i; \boldsymbol{\theta}_t)$$

- $\mathbb{E}[\nabla f_{\mathcal{B}}(\boldsymbol{\theta})] = \nabla f(\boldsymbol{\theta})$
- $\mathcal{O}(1/\sqrt{T}) \quad \mathcal{O}(1/T)$
- batch size

SGD      ill-conditioned

## SGD with Momentum (Polyak)

$$\begin{aligned}\mathbf{v}_t &= \beta \mathbf{v}_{t-1} + \eta \nabla f(\boldsymbol{\theta}_t) \\ \boldsymbol{\theta}_{t+1} &= \boldsymbol{\theta}_t - \mathbf{v}_t\end{aligned}$$

## Nesterov Accelerated Gradient (NAG)

$$\begin{aligned}\mathbf{v}_t &= \beta \mathbf{v}_{t-1} + \eta \nabla f(\boldsymbol{\theta}_t - \beta \mathbf{v}_{t-1}) \\ \boldsymbol{\theta}_{t+1} &= \boldsymbol{\theta}_t - \mathbf{v}_t\end{aligned}$$



■ NAG "Look-ahead"

■  $\beta = 0.9$

# Adam (Adaptive Moment Estimation)

$$\mathbf{m}_t = \beta_1 \mathbf{m}_{t-1} + (1 - \beta_1) \mathbf{g}_t$$

$$\mathbf{v}_t = \beta_2 \mathbf{v}_{t-1} + (1 - \beta_2) \mathbf{g}_t^2$$

$$\hat{\mathbf{m}}_t = \mathbf{m}_t / (1 - \beta_1^t), \quad \hat{\mathbf{v}}_t = \mathbf{v}_t / (1 - \beta_2^t)$$

$$\boldsymbol{\theta}_{t+1} = \boldsymbol{\theta}_t - \eta \cdot \frac{\hat{\mathbf{m}}_t}{\sqrt{\hat{\mathbf{v}}_t} + \epsilon}$$

- $\hat{\mathbf{m}}_t$
- $\hat{\mathbf{v}}_t$     RMS
- $\beta_1 = 0.9, \beta_2 = 0.999, \epsilon = 10^{-8}$
- AdamW    weight decay   AdaBelief   LAMB

# (Learning Rate Schedule)

$\eta$

■ **Step decay**      epoch

■ **Cosine annealing**

$$\eta_t = \eta_{\min} + \frac{1}{2}(\eta_{\max} - \eta_{\min}) \left(1 + \cos \frac{t\pi}{T}\right)$$

■ **Warmup**      epoch       $\eta$  Transformer

■ **Cyclic / OneCycle**      LR

## Warmup

Transformer      Warmup      warmup      Adam

- **Natural Gradient / K-FAC** Fisher Hessian
- **Shampoo** Google
- **AdaHessian** Hutchinson Hessian
- **Newton Sketch** Newton

Hessian  $n \times n$   $\rightarrow$  Hessian

> 100B GPU

**(Data Parallel) + All-Reduce**

- GPU mini-batch
- All-Reduce

**(Model Parallel)**

- Tensor Parallelism GPU
- Pipeline Parallelism GPU

- **Gradient Accumulation** batch

- **(FP16/BF16)**

- **ZeRO** (Zero Redundancy Optimizer)

# ZeRO (DeepSpeed)

Microsoft DeepSpeed ZeRO GPU

|        |                |     |
|--------|----------------|-----|
|        |                |     |
| ZeRO-1 | (Adam $m, v$ ) | 4×  |
| ZeRO-2 | +              | 8×  |
| ZeRO-3 | + +            | GPU |

175B GPT-3 GPU

(Full Fine-tuning)

- 
- 16-32 + +

**(PEFT)**

- **LoRA**  $\Delta W = BA$
- **Adapter**
- **Prefix / Prompt Tuning** token

—



RLHF (Reinforcement Learning from Human Feedback)

- 1 SFT (Supervised Fine-Tuning)
- 2 Reward Model Bradley-Terry
- 3 PPO (Proximal Policy Optimization) Reward Model

$$\max_{\theta} \mathbb{E} [r(x, y) - \beta \cdot \text{KL}(\pi_{\theta} \parallel \pi_{\text{ref}})]$$

**DPO (Direct Preference Optimization)** — RL                      reward                      PPO

$$\mathcal{L}_{\text{DPO}}(\theta) = -\mathbb{E} \left[ \log \sigma \left( \beta \log \frac{\pi_{\theta}(y_w|x)}{\pi_{\text{ref}}(y_w|x)} - \beta \log \frac{\pi_{\theta}(y_l|x)}{\pi_{\text{ref}}(y_l|x)} \right) \right]$$

- **(GPTQ, AWQ)** FP16 4-bit 8-bit
  - (OBS)
  - $\min_{\hat{W}} \|WX - \hat{W}X\|_F^2$
- **KV Cache** Key-Value
- **(Speculative Decoding)** token
- **FlashAttention** IO-aware

Pareto

## LLM

- **(Loss Spike)**
- FP16/BF16     underflow / overflow
- **(Gradient Clipping)**  $\|\mathbf{g}\| > \tau \quad \mathbf{g} \leftarrow \tau \cdot \mathbf{g} / \|\mathbf{g}\|$
- **Z-loss / Logit Regularization**     logit
- output embedding     LayerNorm

**Chinchilla Scaling Law**      $C$       $N$      token      $D$

$$N_{\text{opt}} \propto C^{0.5}, \quad D_{\text{opt}} \propto C^{0.5}$$

- **Decision-Focused Learning** ML
  - $\rightarrow$  LP/IP
- **Learning to Optimize** RL
  - Pointer Network / Transformer TSP, VRP
  - + ML GNN MILP
- **Convex ML** + Elastic Net
- **Bilevel Optimization**

# GPU

ML — GPU SIMT

**CPU ( )**

- 8–64
- 
- 

**GPU ( )**

- 
- (SIMT)
- ( $\sim 2$  TB/s)

## NVIDIA CUDA

- Thread  $\rightarrow$  Warp (32 threads)  $\rightarrow$  Block  $\rightarrow$  Grid
- Global  $\rightarrow$  L2  $\rightarrow$  Shared Memory  $\rightarrow$  Registers
- Tensor Cores (Volta+)  $4 \times 4$

A100 80GB: 6912 CUDA cores + 432 Tensor Cores, 312 TFLOPS (FP16)

# GPU

(SGD, Adam, AdamW)

- 1 Forward pass  $\mathbf{y} = \mathbf{W}\mathbf{x}$
- 2 Loss
- 3 Backward pass  $\nabla_{\mathbf{W}}\mathcal{L} = (\nabla_{\mathbf{y}}\mathcal{L})\mathbf{x}^{\top}$
- 4

1 3 GPU 50×–200×

- cuBLAS GEMM ( )
- cuDNN GPU kernel
- (Newton, K-FAC) Hessian GPU

GPU

175B ~350 GB (FP16) TB

## (Mixed Precision)

FP32          FP16

(NVIDIA APEX / PyTorch AMP)

- 1      FP32
- 2      Forward/Backward      FP16      2×
- 3      **Loss Scaling**           $S$       underflow
- 4      FP32

### BF16 (Brain Float 16)

- FP32      (8 bit)      (7 bit)
- = FP32      loss scaling
- NVIDIA A100 / H100

Tensor Cores      FP16/BF16      FP32      8–16×

GPU batch size

### (Gradient Accumulation)

- $k$  micro-batch +
- $k$  batch size  $\times k$
- = micro-batch size  $\approx$
  
- (Linear Scaling Rule, Goyal et al. 2017)  
 $\eta_{\text{new}} = \eta_{\text{base}} \times (B_{\text{new}}/B_{\text{base}})$  warmup
- LARS (Layer-wise Adaptive Rate Scaling)
- LAMB Adam BERT



## (Data Parallel)

- GPU
  - GPU mini-batch
  - All-Reduce
  - PyTorch DDP, Horovod
- 
- Tensor Parallel (Megatron-LM)
  - Pipeline Parallel (GPipe)
  - 3D Parallel = DP + TP + PP

## NCCL

- All-Reduce
- All-Gather
- Reduce-Scatter
- ~600 GB/s (NVLink)

## FSDP (Fully Sharded DP)

+ + ZeRO-3 PyTorch

1B + Adam + FP16

| (FP16)          | 2 GB | ~20% |
|-----------------|------|------|
| Adam $m$ (FP32) | 4 GB | ~40% |
| Adam $v$ (FP32) | 4 GB | ~40% |
| (FP16)          | 2 GB | ~20% |
| (Activations)   |      |      |

■ **Gradient Checkpointing**

■ **Activation Offloading** CPU

■ **FlashAttention** IO-aware + softmax  $\mathcal{O}(N^2)$   $\mathcal{O}(N)$

■ **Kernel Fusion** CUDA kernel Bias + GeLU + Dropout



# GPU Checklist

- 1 AMP (FP16) BF16 Tensor Core
- 2 DataLoader prefetch pin\_memory
- 3 batch + batch
- 4 DDP FSDP / DeepSpeed
- 5 torch.compile / TensorRT XLA kernel fusion
- 6 gradient checkpointing FlashAttention fused optimizer
- 7 nvidia-smi, Nsight Systems (timeline), Nsight Compute (kernel)
- 8 **batch size** GPU ( > 80% )

GPU —

# GPU

- **LP** GPU cuSOLVER
- MIP B&B GPU
- Transformer/GNN GPU TSP/VRP
- GPU
- **RL for OR** GPU PPO /
- **SDP** GPU SDP Cholesky

NVIDIA cuOpt

GPU

VRP

100×

- 1
- 2 KKT /
- 3 ADMM
- 4 **ML** SGD → Momentum → Adam
- 5 **LLM** ZeRO PEFT LoRA RLHF
- 6 ML OR
- 7 **GPU** CUDA FlashAttention cuOpt